

LSTM(Long Short Term Memory)

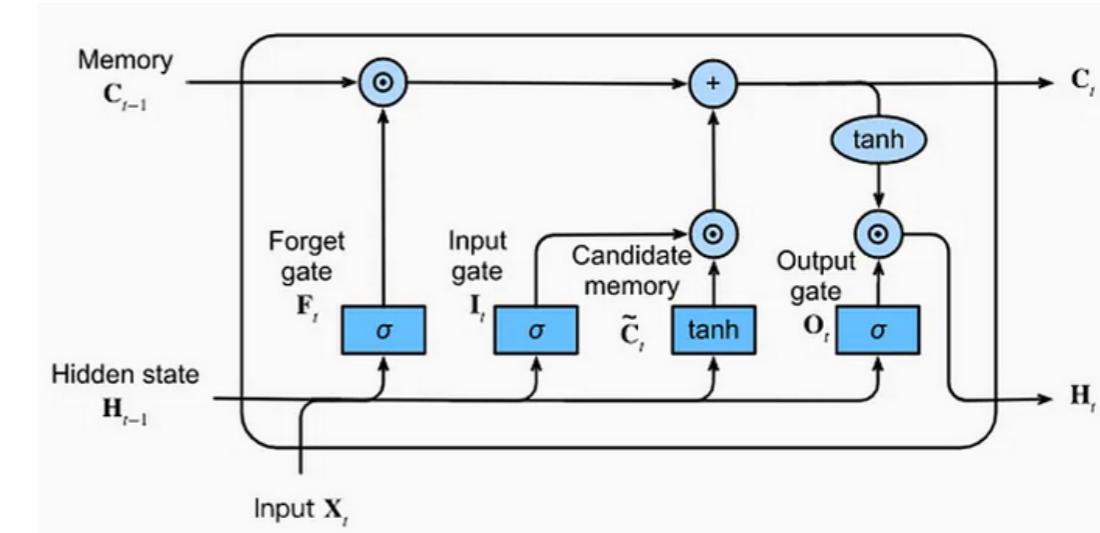
Why LSTM ? Not Rnn

A standard RNN updates its hidden state as:

$$h_t = \tanh(W_h h_{t-1} + W_x x_t)$$

1. During backpropagation through time(BPTT):
Gradients are repeatedly multiplied by the same weight matrix.

2. This causes vanishing gradients (values $\rightarrow 0$) or exploding gradients (values $\rightarrow \infty$).



Practical consequence:

- 1.RNNs forget information quickly
- 2.They struggle when important context appears many time steps earlier.

LSTM performs better than RNN on:

- 1.Language modeling
- 2.Machine translation
- 3.Speech recognition
- 4.Time-series forecasting
- 5.POS tagging and NER

This is why almost all classical NLP pipelines replaced RNNs with LSTMs before Transformers.

Why tanh is stable activation function for rnn and lstm?

tanh is considered stable because it is bounded, zero-centered, has a smooth derivative, and a maximum gradient of 1, making it well-suited for repeated application across time steps in recurrent architectures. It avoids the exploding activations of ReLU and the biased gradients of sigmoid, while working harmoniously with gating mechanisms in LSTMs to preserve long-term dependencies.

tanh does not eliminate vanishing gradients by itself.

What eliminates vanishing gradients in LSTM is:

$$c_t = f_t \odot c_{t-1} + i_t \odot \tilde{c}_t$$

tanh simply does not make the problem worse.

Why LSTM uses tanh specifically?

In LSTM:

- 1.sigmoid \rightarrow control gates (how much to allow)
- 2.tanh \rightarrow content values (what to store/output)

Example:

$$c_t = \tanh(W_c[h_{t-1}, x_t])$$

Why this is ideal:

- 1.Keeps memory values bounded
- 2Allows both positive and negative information
- 3.Stable gradient flow into cell state

LSTM Arcitecture and gates

Forget Gate:

Decides how much of the old memory to keep.

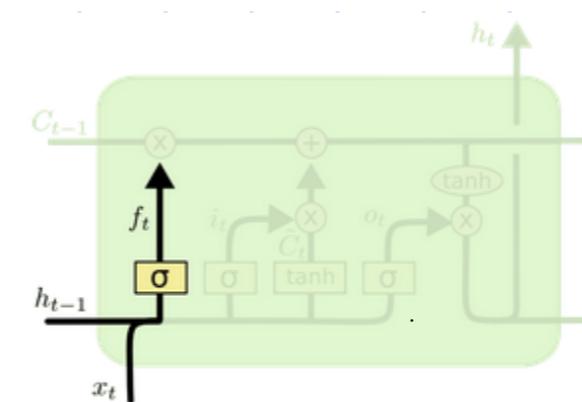
$$f_t = \sigma(W_f \cdot [h_{t-1}, x_t] + b_f)$$

Output range: [0, 1]

0 → forget everything

1 → keep everything

if earlier information is no longer relevant, the forget gate removes it.



$$f_t = \sigma(W_f \cdot [h_{t-1}, x_t] + b_f)$$

Input Gate:

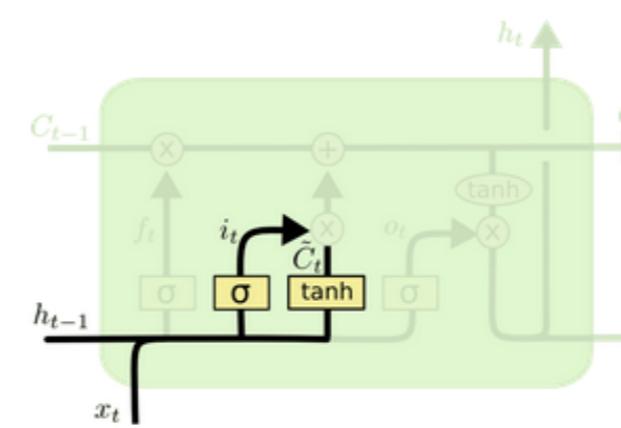
4. Input Gate — “What should I write?”

Step 1: Decide where to write

$$i_t = \sigma(W_i \cdot [h_{t-1}, x_t] + b_i)$$

Step 2: Decide what to write

$$c_t \sim \tanh(W_c \cdot [h_{t-1}, x_t] + b_c)$$



$$i_t = \sigma(W_i \cdot [h_{t-1}, x_t] + b_i)$$

$$\tilde{C}_t = \tanh(W_C \cdot [h_{t-1}, x_t] + b_C)$$

Cell State Update:

This is the most important equation.

$$c_t = f_t \odot c_{t-1} + i_t \odot \tilde{C}_t$$

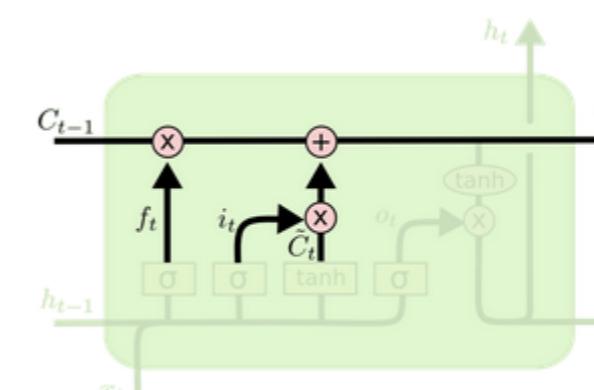
What happens here:

Old memory is scaled, not overwritten

New memory is selectively added

Gradient can flow through time almost unchanged

This is why LSTM solves vanishing gradients.



$$C_t = f_t * C_{t-1} + i_t * \tilde{C}_t$$

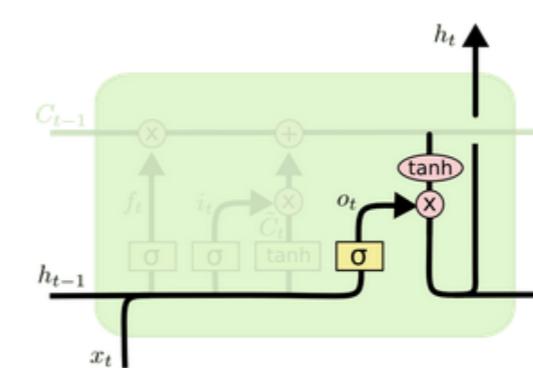
Output Gate:

Step 1: Decide how much to expose

$$o_t = \sigma(W_o \cdot [h_{t-1}, x_t] + b_o)$$

Step 2: Generate hidden state

$$h_t = o_t \odot \tanh(c_t)$$



$$o_t = \sigma(W_o \cdot [h_{t-1}, x_t] + b_o)$$

$$h_t = o_t * \tanh(C_t)$$

Key idea:

Cell state = internal memory

Hidden state = controlled view of that memory

One complete LSTM step (ordered)

At time step t:

Concatenate $[h_{t-1}, x_t]$

Compute forget gate f_t

Compute input gate i_t

Compute candidate memory \tilde{C}_t

Update cell state c_t

Compute output gate o_t

Compute hidden state h_t